

2

---

3 **Attribute-based Access Control for**  
4 **Microservices-based Applications**  
5 **Using a Service Mesh**

---

6

7

8

9

Ramaswamy Chandramouli  
Zack Butcher  
Aradhna Chetal

10

11

12

13

14

15

16

This publication is available free of charge from:  
<https://doi.org/10.6028/NIST.SP.800-204B-draft>

17 **Draft NIST Special Publication 800-204B**

18

19 **Attribute-based Access Control for**  
20 **Microservices-based Applications**  
21 **Using a Service Mesh**

22

23

24

25

Ramaswamy Chandramouli  
*Computer Security Division*  
*Information Technology Laboratory*

26

27

28

Zack Butcher  
*Tetrade*  
*San Francisco, CA*

29

30

31

32

Aradhna Chetal  
*TIAA*  
*New York, NY*

33

34

This publication is available free of charge from:  
<https://doi.org/10.6028/NIST.SP.800-204B-draft>

35

36

37

January 2021



38

39

40

41

42

43

44

45

U.S. Department of Commerce  
*Wynn Coggins, Acting Secretary*

National Institute of Standards and Technology  
*James K. Olthoff, Performing the Non-Exclusive Functions and Duties of the Under Secretary of Commerce*  
*for Standards and Technology & Director, National Institute of Standards and Technology*

46

## Authority

47 This publication has been developed by NIST in accordance with its statutory responsibilities under the  
48 Federal Information Security Modernization Act (FISMA) of 2014, 44 U.S.C. § 3551 *et seq.*, Public Law  
49 (P.L.) 113-283. NIST is responsible for developing information security standards and guidelines, including  
50 minimum requirements for federal information systems, but such standards and guidelines shall not apply  
51 to national security systems without the express approval of appropriate federal officials exercising policy  
52 authority over such systems. This guideline is consistent with the requirements of the Office of Management  
53 and Budget (OMB) Circular A-130.

54 Nothing in this publication should be taken to contradict the standards and guidelines made mandatory and  
55 binding on federal agencies by the Secretary of Commerce under statutory authority. Nor should these  
56 guidelines be interpreted as altering or superseding the existing authorities of the Secretary of Commerce,  
57 Director of the OMB, or any other federal official. This publication may be used by nongovernmental  
58 organizations on a voluntary basis and is not subject to copyright in the United States. Attribution would,  
59 however, be appreciated by NIST.

60 National Institute of Standards and Technology Special Publication 800-204B  
61 Natl. Inst. Stand. Technol. Spec. Publ. 800-204B, 37 pages (January 2021)  
62 CODEN: NSPUE2

63 This publication is available free of charge from:  
64 <https://doi.org/10.6028/NIST.SP.800-204B-draft>

65 Certain commercial entities, equipment, or materials may be identified in this document in order to describe an  
66 experimental procedure or concept adequately. Such identification is not intended to imply recommendation or  
67 endorsement by NIST, nor is it intended to imply that the entities, materials, or equipment are necessarily the best  
68 available for the purpose.

69 There may be references in this publication to other publications currently under development by NIST in accordance  
70 with its assigned statutory responsibilities. The information in this publication, including concepts and methodologies,  
71 may be used by federal agencies even before the completion of such companion publications. Thus, until each  
72 publication is completed, current requirements, guidelines, and procedures, where they exist, remain operative. For  
73 planning and transition purposes, federal agencies may wish to closely follow the development of these new  
74 publications by NIST.

75 Organizations are encouraged to review all draft publications during public comment periods and provide feedback to  
76 NIST. All NIST Computer Security Division publications, other than the ones noted above, are available at  
77 <http://csrc.nist.gov/publications>.

78 **Public comment period: *January 27, 2021 through February 24, 2021***

79 National Institute of Standards and Technology  
80 Attn: Computer Security Division, Information Technology Laboratory  
81 100 Bureau Drive (Mail Stop 8930) Gaithersburg, MD 20899-8930  
82 Email: [sp800-204b-comments@nist.gov](mailto:sp800-204b-comments@nist.gov)

83 All comments are subject to release under the Freedom of Information Act (FOIA).  
84

85

## Reports on Computer Systems Technology

86 The Information Technology Laboratory (ITL) at the National Institute of Standards and  
87 Technology (NIST) promotes the U.S. economy and public welfare by providing technical  
88 leadership for the Nation's measurement and standards infrastructure. ITL develops tests, test  
89 methods, reference data, proof of concept implementations, and technical analyses to advance the  
90 development and productive use of information technology. ITL's responsibilities include the  
91 development of management, administrative, technical, and physical standards and guidelines for  
92 the cost-effective security and privacy of other than national security-related information in federal  
93 information systems.

94

### Abstract

95 Deployment architecture in cloud-native applications now consists of loosely coupled  
96 components, called microservices, with all application services provided through a dedicated  
97 infrastructure, called service mesh, independent of the application code. Two critical security  
98 requirements in this architecture are (a) to build the concept of zero trust by enabling mutual  
99 authentication in communication between any pair of services and (b) a robust access control  
100 mechanism based on an access control such as ABAC that can be used to express a wide set of  
101 policies and is scalable in terms of user base, objects (resources), and deployment environment.  
102 This document provides deployment guidance for building an authentication and authorization  
103 framework within the service mesh that meets these requirements. A reference platform for  
104 hosting the microservices-based application and a reference platform for the service mesh are  
105 included to illustrate the concepts in the recommendations and provide the context in terms of  
106 the components used in real-world deployments.  
107

108

### Keywords

109 attribute-based access control; authentication policy; authorization policy; CI/CD; DevSecOps;  
110 JSON web token; microservices-based application; mutual TLS; next generation access control;  
111 policy enforcement point; role-based access control; service mesh; service proxy; zero trust.

112

113

114

## Acknowledgments

115 The authors like to express their sincere thanks to Mr. David Ferraiolo of NIST for initiating this  
116 effort to provide a targeted deployment guidance in the form of an authentication and  
117 authorization framework in a service mesh environment used for protecting microservices-based  
118 applications. They also express thanks to Isabel Van Wyk of NIST for her detailed editorial  
119 review.

120

## Call for Patent Claims

121 This public review includes a call for information on essential patent claims (claims whose use  
122 would be required for compliance with the guidance or requirements in this Information  
123 Technology Laboratory (ITL) draft publication). Such guidance and/or requirements may be  
124 directly stated in this ITL Publication or by reference to another publication. This call also  
125 includes disclosure, where known, of the existence of pending U.S. or foreign patent applications  
126 relating to this ITL draft publication and of any relevant unexpired U.S. or foreign patents.

127 ITL may require from the patent holder, or a party authorized to make assurances on its behalf,  
128 in written or electronic form, either:

129 a) assurance in the form of a general disclaimer to the effect that such party does not hold  
130 and does not currently intend holding any essential patent claim(s); or

131 b) assurance that a license to such essential patent claim(s) will be made available to  
132 applicants desiring to utilize the license for the purpose of complying with the guidance  
133 or requirements in this ITL draft publication either:

134 i. under reasonable terms and conditions that are demonstrably free of any unfair  
135 discrimination; or

136 ii. without compensation and under reasonable terms and conditions that are  
137 demonstrably free of any unfair discrimination.

138 Such assurance shall indicate that the patent holder (or third party authorized to make assurances  
139 on its behalf) will include in any documents transferring ownership of patents subject to the  
140 assurance, provisions sufficient to ensure that the commitments in the assurance are binding on  
141 the transferee, and that the transferee will similarly include appropriate provisions in the event of  
142 future transfers with the goal of binding each successor-in-interest.

143 The assurance shall also indicate that it is intended to be binding on successors-in-interest  
144 regardless of whether such provisions are included in the relevant transfer documents.

145 Such statements should be addressed to: [sp800-204b-comments@nist.gov](mailto:sp800-204b-comments@nist.gov)

146

## 147 **Executive Summary**

148 Two significant features of the application environment in emerging cloud-native applications  
149 are:

- 150 • Applications have multiple loosely coupled components called microservices that  
151 communicate with each other across the network.
- 152 • A dedicated infrastructure called the service mesh provides all services for the application  
153 (e.g., authentication, authorization, routing, network resilience, security monitoring),  
154 which can be deployed independently of the application code.

155 With the disappearance of a network perimeter because of the need to provide ubiquitous access  
156 to applications from multiple remote locations using different types of devices, it is necessary to  
157 build the concept of zero trust into the application environment. Further, the cloud-native  
158 applications span different domains and, therefore, require increased precision in specifying  
159 policy by considering a large set of variables. The service mesh provides a framework for  
160 building these and other operational assurances.

161  
162 The framework includes:

- 163 • An authenticatable runtime identity for services, the ability to authenticate application  
164 (user) credentials, and encryption of communication in transit and between services
- 165 • A Policy Enforcement Point (PEP) that is separately deployable and controllable from the  
166 application; the service mesh's side-car proxies
- 167 • Logs and metrics for monitoring policy enforcement

168  
169 The service mesh's native feature to authenticate end-user credentials attached to the request  
170 (e.g., using a Java Web Token [JWT]) is augmented in many offerings to provide the ability to  
171 call external authentication and authorization systems on behalf of the application. The capability  
172 to deploy these authentication and authorization systems as services in the mesh also provides  
173 operational assurances for encryption in transit, identity, a PEP, authentication, and authorization  
174 for end-user identity.

175  
176 The objective of this document is to provide deployment guidance for an authentication and  
177 authorization framework within a service mesh for microservices-based applications that  
178 leverages the features listed above. A reference platform for hosting the microservices-based  
179 application and the service mesh is included to illustrate the concepts in the recommendations  
180 and provide context in terms of the components used in real-world deployments.

181

182  
183  
184  
185  
186  
187  
188  
189  
190  
191  
192  
193  
194  
195  
196  
197  
198  
199  
200  
201  
202  
203  
204  
205  
206  
207  
208  
209  
210  
211  
212  
213

**Table of Contents**

**Executive Summary ..... iv**

**1 Introduction ..... 1**

    1.1 Service Mesh Capabilities..... 1

    1.2 Candidate Applications ..... 3

    1.3 Scope..... 3

    1.4 Target Audience..... 3

    1.5 Relationship to other NIST Guidance Documents ..... 3

    1.6 Organization of this document ..... 4

**2 Microservices-based Application and Service Mesh – Reference Platforms ... 5**

    2.1 Reference Platform for Hosting a Microservices-based Application ..... 5

        2.1.1 Limitations of Reference Hosting Platform for Security ..... 5

    2.2 Service Mesh Reference Platform – Conceptual Architecture ..... 6

        2.2.1 Service Mesh Functions for Reference Hosting Platform ..... 7

**3 Attribute-based Access Control (ABAC) – Background ..... 9**

    3.1 ABAC Deployment for Microservices-based Applications Using Service Mesh  
        12

**4 Authentication and Authorization Policy Configuration in Service Mesh..... 13**

    4.1 Hosting Platform Configuration ..... 13

    4.2 Service Mesh Configuration..... 13

    4.3 Higher-level Security Configuration Parameters ..... 14

    4.4 Authentication Policies..... 15

        4.4.1 Specifying Authentication Policies..... 16

        4.4.2 Service-level Authentication ..... 16

        4.4.3 End User Authentication..... 17

    4.5 Authorization Policies..... 17

        4.5.1 Service-level Authorization Policies..... 18

        4.5.2 End-user Level Authorization Policies ..... 18

        4.5.3 Model-based Authorization Policies..... 20

    4.6 Authorization Policy Elements..... 21

        4.6.1 Policy Types..... 21

        4.6.2 Policy Target or Authorization Scope ..... 21

214            4.6.3 Policy Sources ..... 22

215            4.6.4 Policy Operations ..... 22

216            4.6.5 Policy Conditions..... 22

217    **5 ABAC Deployment for Service Mesh..... 24**

218            5.1 Security Assurance for Authorization Framework Enforcement..... 24

219            5.2 Supporting Infrastructure for ABAC Authorization Framework..... 24

220            5.2.1 Service-to-Service Request (SVC-SVC) – Supporting Infrastructure . 24

221            5.2.2 End User + Service-to-Service Request (EU+SVC-SVC) – Supporting

222            Infrastructure ..... 25

223            5.3 Advantages of ABAC Authorization Framework for Service Mesh..... 26

224            5.4 Enforcement Alternatives in Proxies ..... 26

225    **6 Summary and Conclusions ..... 27**

226    **References ..... 28**

227



## 228 **1 Introduction**

229 Applications based on microservices-based architecture and an application infrastructure based  
230 on service mesh that provides various security services through service proxies have emerged as  
231 the widespread application environment for cloud-native applications. With the disappearance of  
232 the network perimeter due to the need to provide ubiquitous access to these applications from  
233 multiple remote locations using different types of devices, it is necessary to build the concept of  
234 zero trust [1] into this application environment. Further, the loosely coupled nature of the  
235 components of these cloud-native applications (i.e., microservices) facilitates independent  
236 design, development, and agile deployment (e.g., CI/CD [2]) of the constituent microservices,  
237 enabling paradigms such as DevSecOps [3] need to be used.

238 The security requirements for microservices-based applications are discussed extensively in [4]  
239 and summarized here to provide context for this discussion. They are:

- 240 ● Multiple, loosely coupled microservices communicate through network calls, and these  
241 communication links must be protected. In the case of monolithic applications, these  
242 communications take place through procedure calls.
- 243 ● The entire network is untrusted, and each microservice is untrusted. Therefore, mutual  
244 authentication between microservices and secure communication channels between  
245 paired microservices through mechanisms such as mutual TLS (mTLS) are required.
- 246 ● The logging data that pertains to each microservice must be consolidated to obtain a  
247 security profile in order for forensics, audits, and analytics to assess the overall health of  
248 the application.

249 Operating in multiple security domains and multiple clouds, cloud-native applications require a  
250 secure authentication and authorization framework. The critical requirements of this framework  
251 when implemented within the service mesh are:

- 252 ● The code that is part of this framework is verifiable and non-bypassable (always  
253 invoked), thus satisfying the requirements of a security kernel.
- 254 ● The framework should provide authentication and authorization services at multiple  
255 levels: service and end-user.
- 256 ● The framework should be able to support a diverse set of authorization policies.

257 The operational assurances required for meeting the above requirements and others are provided  
258 by the service mesh. The specific features in service that enable these are given in the next  
259 section.

### 260 **1.1 Service Mesh Capabilities**

261 A service mesh provides a framework for building a set of operational assurances for an  
262 organization. That framework includes an authenticatable runtime identity for services, the  
263 ability to authenticate application (user) credentials, encryption in transit of communication  
264 between services, a Policy Enforcement Point (PEP) separately deployable and controllable from

265 the application (the service mesh’s side-car proxies), and logs and metrics for monitoring policy  
266 enforcement. Using these mesh features, a set of controls can be built for all applications that are  
267 part of the mesh (e.g., all traffic is encrypted, all traffic to an application goes through the side-  
268 car [PEP]). These controls provide a set of operational assurances for applications in an  
269 organization deployed in the service mesh.

270 A significant benefit of the service mesh architecture is that the key piece that allows for these  
271 controls to be built—the sidecar proxy deployed next to every application—has more security  
272 benefits than the traditional approach of building these operational assurances into the  
273 application code. First, the life cycle of the sidecar is independent of the application, making it  
274 easier to manage across a fleet (e.g., push updates, ensure a consistent version is deployed  
275 everywhere). Second, modern implementations (like Istio) allow for dynamic configuration. It is  
276 easy to update policies, and updates take effect immediately and without having to redeploy  
277 applications. Finally, the mesh’s centralized control allows security teams to build policies that  
278 apply to the entire organization so that application developers who build business value are  
279 secure by default.

280 A service mesh provides the ability to authenticate end user credentials attached to the request,  
281 like a JSON Web Token (JWT). Many service meshes (e.g., Istio) go further and provide the  
282 ability for the mesh’s sidecar to call external authentication and authorization systems on behalf  
283 of the application. This grants the ability to move request-level policy enforcement out of the  
284 application code, trusting instead on the mesh’s assurance that requests that reach the service  
285 have been authenticated and authorized for the action that the request is taking. The mesh can  
286 even be configured to pass proof of this to the application. This, coupled with the service mesh’s  
287 centralized control, means it is possible for a central team to mandate and manage application-  
288 level security across the entire organization, delegating to individual application teams only to  
289 specify what permissions are required for each applications’ actions.

290 Using the service mesh architecture also means that authentication and authorization systems can  
291 be deployed as services in the mesh. Like any other service in the mesh, they benefit from the  
292 operational assurances the mesh provides: encryption in transit, identity, a PEP, authentication,  
293 and authorization for end user identity. This makes it cheaper to operate an organization’s  
294 authentication and authorization systems securely and reliably.

295 In addition to the service mesh features, the capabilities of the access control model play an  
296 important role in the authentication and authorization framework. Attribute-based access control  
297 (ABAC) has emerged as a promising approach for supporting multiple authorization policies  
298 (third requirement above). As per [5], ABAC is defined as “an access control method where  
299 subject requests to perform operations on objects are granted or denied based on assigned  
300 attributes of the subject, assigned attributes of the object, (optionally) environmental conditions,  
301 and a set of policies that are specified in terms of those attributes and conditions.” The main  
302 focus of this document is to provide guidance on an authentication and authorization framework,  
303 the latter using ABAC to secure microservices-based applications using service mesh.

## 304 1.2 Candidate Applications

305 The service mesh is most widely used today with containerized applications but can be extended  
306 into other environments, such as stateful applications.

## 307 1.3 Scope

308 This document focuses on providing guidance for building a secure authentication and  
309 authorization framework using components of a service mesh for securing services in  
310 microservice-based applications. A reference application hosting platform and a reference  
311 service mesh platform have been used as examples to illustrate the recommendations in the  
312 context of real-world application artifacts (e.g., containers, VMs, etc.). The chosen reference  
313 application platform is the open-source Kubernetes, and the chosen reference service mesh  
314 platform is Istio. Application infrastructure components in the service mesh that provide other  
315 services like network routing, network resilience, and monitoring are outside of the scope of this  
316 document.

## 317 1.4 Target Audience

318 The target audience of the guidance document for developing an authentication and authorization  
319 framework for microservices-based applications using the service mesh includes:

- 320 ● Security solutions architects who want to protect the application workloads in microservices-  
321 based applications.
- 322 ● Platform architects who want to incorporate a service mesh into the platform offered by their  
323 organization to its developers
- 324 ● Developers who want to develop authentication and authorization plug-ins in this application  
325 environment

## 326 1.5 Relationship to other NIST Guidance Documents

327 This guidance document focuses on building an authentication and authorization framework  
328 within the service mesh used for securing microservices-based applications. The following  
329 publications provide background information for the contents of this document:

- 330 ● Special Publication (SP) 800-204, *Security Strategies for Microservices-based Application*  
331 *Systems* [4], discusses the characteristics of microservices-based applications and the overall  
332 security requirements and strategies for addressing those requirements.
- 333 ● Special Publication (SP) 800-204A, *Building Secure Microservices-based Applications*  
334 *Using Service-Mesh Architecture* [6], provides deployment guidance for various security  
335 services (e.g., authentication and authorization, security monitoring, etc.) for a microservices-  
336 based application using a dedicated infrastructure (i.e., a service mesh) that uses service  
337 proxies that operate independent of the application code.

**338 1.6 Organization of this document**

339 The organization of this document is as follows:

- 340 ● Chapter 2 provides an overview of a microservices-based application, its security  
341 requirements, components of a service mesh, and a brief description of the overall  
342 architecture of the reference hosting platform and the reference service mesh platform.  
343 The latter two are used as examples to illustrate the building blocks involved in the  
344 deployment recommendations.
- 345 ● Chapter 3 outlines the advantages of ABAC for the application environment and  
346 describes the functional architecture for two of the standard ABAC representations.
- 347 ● Chapter 4 discusses the building blocks of the authentication and authorization  
348 framework, the basic configuration that is required in the reference hosting and reference  
349 service mesh platform for implementing the framework, and the salient features of the  
350 framework.
- 351 ● Chapter 5 provides recommendations regarding deployment of the various use cases  
352 pertaining to authorization policies as well as the building blocks (policy components) of  
353 these policies.
- 354 ● Chapter 6 provides the summary and conclusions.

355

## 2 Microservices-based Application and Service Mesh – Reference Platforms

The objective of this document is to offer recommendations for the deployment of an authentication and authorization framework for microservices-based applications within a service mesh that provides the infrastructure for various services, including critical security services. A reference platform for hosting microservices-based applications and the service mesh is included to provide clarity and context for concepts and recommendations in real-world application environments. A brief description of these reference platforms is also provided in terms of their overall architecture and salient building blocks.

### 2.1 Reference Platform for Hosting a Microservices-based Application

Kubernetes is an orchestration and resource management system widely used for microservices-based applications. In a large application, there will be several microservices, each of which is implemented as a container. Scalable, automated means are required for deployments, operations, upgrading services, and monitoring the health of these containers. The Kubernetes architecture provides the tools to achieve these goals.

To enable application-level, fine-grained access control, it is imperative to have some cluster-level security mechanisms for the clusters that are configured using the hosts of the application components (i.e., microservices). Considering a scenario where the host is a worker node of a Kubernetes platform cluster and the application components are running inside of a container with a pod (i.e., a group of containers) as a deployment artifact, the following cluster-level security measures are required. These measures are defined and enforced through artifacts called pod security policies.

For example, one of the most well-known features of Kubernetes is pod-level *horizontal scaling*. This means that when services receive more traffic, more instances will be generated across machines that grow or shrink on demand. Kubernetes supports auto-vertical scaling on the pod level. Thus, a cluster could be configured to scale the machine on which a pod runs up or down to more accurately fit the anticipated power needs of any microservice. For example, if certain subsets of worker nodes saw spikes in traffic at key times, with the right usage analysis, one could potentially reschedule across machines in order to save costs and optimize performance [7].

Similarly, Kubernetes offers features to monitor the health of the microservices (check the status and readiness). The data to perform these functions is configured in declarative deployment documents, typically as YAML, that describe the port that a pod's containers are listening on. One can specify what to do when services do not start, do not perform as normal, or exit unexpectedly.

#### 2.1.1 Limitations of Reference Hosting Platform for Security

Microservices-based applications require several application infrastructure and security services, such as authentication, authorization, monitoring, logging, auditing, traffic control, caching,

393 secure ingress, service-to-service, and egress communication. Moreover, the following  
394 advantages of API architecture are not fully leveraged in the reference platform [8]:

- 395 • A unified way to apply cross-cutting concerns
- 396 • Out of the box plugins to apply cross-cutting concerns quickly
- 397 • A framework for building custom plugins
- 398 • Managing security in a single plane
- 399 • Reduced operation complexity
- 400 • Easy governance of third-party developers and integrators
- 401 • Saving the cost of development and operations

402 By default, communication between Kubernetes containers is insecure, and there is no easy way  
403 to enforce TLS between pods since this would result in individually maintaining hundreds of  
404 TLS certificates. Pods that communicate do not apply identity and access management between  
405 themselves. Though there are tools, such as [Kubernetes Network Policy](#), that can be  
406 implemented to act as a firewall between pods, they are a [layer 3](#) solution rather than a [layer](#)  
407 [7](#) solution, which is what most modern firewalls are. This means that while one can know the  
408 source of traffic, one cannot peek into the data packets to understand what they contain. It does  
409 not allow for making vital metadata-driven decisions, such as routing on a new version of a pod  
410 based on an HTTP header. There are Kubernetes ingress objects that do provide a reverse proxy  
411 based on layer 7, but they do not offer anything more than simple traffic routing. Kubernetes  
412 does offer different ways of deploying pods that do some form of [A/B testing](#) or canary  
413 deployments, but they are done at the connection level and provide no fine-grained control or  
414 fast failback. For example, if a developer wants to deploy a new version of a microservice and  
415 pass 10 % of traffic through it, they will have to scale the containers to at least 10—nine for  
416 the old version and one for the new version. Further, Kubernetes cannot split the traffic  
417 intelligently and instead balances loads between pods in a round-robin fashion. Every  
418 Kubernetes container within a pod has separate log, and a custom solution over Kubernetes must  
419 be implemented to capture and consolidate them.

420 Although the Kubernetes dashboard offers features like monitoring pods and checking their  
421 health, it does not expose metrics that describe how application components interact with each  
422 other, how much traffic flows through each of the pods, or what chains of containers make up the  
423 application. Since traffic flow cannot be traced through Kubernetes pods out of the box, it is  
424 unclear where on the chain the failure for the request occurred.

425 A service mesh addresses these limitations [9]. This document will first consider the service  
426 mesh architecture, followed by implementation of service mesh capabilities in the context of the  
427 reference platform (Kubernetes).

## 428 **2.2 Service Mesh Reference Platform – Conceptual Architecture**

429 A service mesh is the network of microservices that make up applications and the interactions  
430 between them. It helps to manage microservices-based applications using two major  
431 components:

- 432 1. **Data Plane.** This is the component that performs the actual routing or communication of  
433 messages between microservices. It also gathers telemetry data, which helps to monitor  
434 the health and state of the services. The traffic that flows through the data plane is thus  
435 the application-related (business) data.
- 436 2. **Control Plane.** This is the component that provides an API to define policies. This API  
437 is often independent of the platform on which the microservices application and, hence,  
438 the Service Mesh runs. The control plane also helps the administrator populate the data  
439 plane component with configuration that determines how to route traffic. The control  
440 plane is the brain of a service mesh. The traffic that flows through the control plane  
441 consists of messages of interaction between service mesh components.

442 The control plane may consist of multiple modules, and the distribution of functionality among  
443 these modules may be different in different service mesh offerings. However, they all provide the  
444 following core functions:

- 445 a. A module that parses the policy rules defined in the control plane and converts them into  
446 configuration parameters in the data plane module (i.e., the sidecar proxy). These policies  
447 may pertain to various functions, such as authentication and authorization, service  
448 discovery, traffic management (including load balancing), intelligent routing, blue-green  
449 deployments, canary rollouts, and much more. It may also include configuration  
450 parameters related to resiliency in the service mesh, such as timeout, retry, and circuit-  
451 breaking capabilities.
- 452 b. A module that provides all of the infrastructure functionality for authentication,  
453 authorization, and establishing a secure, encrypted session while two microservices  
454 communicate. These functions include user authentication, credential management,  
455 digital certificate management, and traffic encryption.

### 456 2.2.1 Service Mesh Functions for Reference Hosting Platform

457 In order to describe the generic service mesh functions in the context of the reference platform—  
458 which, in this case, is Kubernetes—the deployment details of both the microservices application  
459 and service mesh components in that platform must be considered. Since authentication and  
460 authorization functions are the focus of this document, discussions for those functions on the  
461 Kubernetes platform will be confined to the functions in the service mesh.

462 Since the sidecar proxy code implemented as a container is hosted in the same pod as the  
463 microservice container, they share the same network namespace and are present in the same node  
464 (e.g., VM or a physical machine). Both containers have the same IP address and share the same  
465 IP Table rules. That makes the proxy take complete control over the pod and handle all traffic  
466 that passes through it [10].

467 Taking the example of establishing a mutual TLS session, the proxy will interact with the  
468 module in the control plane of the service mesh to check whether it needs to encrypt traffic  
469 through the chain and establish mutual TLS with the backend pod. Enabling this functionality  
470 using mutual TLS requires every pod to have a certificate (i.e., a valid credential), and since a

471 good-sized microservice application may be hosted in hundreds of pods, this may involve  
472 managing hundreds of short-lived certificates. This in turn requires the service mesh to have a  
473 robust identity, access manager, certificate store, and certificate validation. In addition,  
474 mechanisms for identifying and authenticating the two communicating pods are required for  
475 supporting authentication policies.

476 A service mesh not only provides various application services during runtime but also supports  
477 the DevSecOps development and maintenance paradigm. The development team can concentrate  
478 their efforts on efficient development paradigms, such as code modularity and structuring,  
479 without worrying about the security and management details of their implementation.

480 The service mesh is reference platform-aware and thus automatically injects sidecar containers  
481 into the pods. Once the service mesh inserts the sidecar containers, operations and security teams  
482 can enforce policies on the traffic and help secure and operate the application. These teams can  
483 also configure monitoring of the microservices applications without interfering with the  
484 functioning of the applications.

485



### 3 Attribute-based Access Control (ABAC) – Background

487 Attribute-based access control (ABAC) is an authorization framework or engine that computes  
488 decisions for user access requests based on attributes and policies expressed in terms of attributes  
489 [5]. The advantages of ABAC for microservices-based applications using service mesh include:

- 490 • Cloud-native applications span different domains and require increased precision in  
491 specifying policy by considering a large set of variables. Because of its scalability with  
492 respect to attribute and value stores and associated policies, ABAC can meet this  
493 requirement.
- 494 • Attributes and their values associated with users, application objects, resources, and  
495 environments are independently assigned. Hence, policies based on the attributes do not  
496 create a tight subject/object relationship since access decisions are ultimately dependent on  
497 dynamic attribute values.
- 498 • Policies are expressed in terms of attributes without prior knowledge of potentially numerous  
499 users and resources that are or will be governed under those policies, and users and resources  
500 are independently assigned attribute values without knowledge of policy details. This dual  
501 feature enables access control decisions to be based on centralized, enterprise-wide policies  
502 while also supporting the DevSecOps approach that provides autonomy to each microservice  
503 development team to make all decisions regarding their module, including the resource  
504 attribute assignments.

505 Due to the features described above, the ABAC authorization framework is a natural fit for the  
506 class of cloud-native applications whose design is based on splitting an application into several  
507 loosely coupled modules called microservices with each being developed and deployed by  
508 independent teams.

509 The ABAC framework has two standardized, representational structures. One uses a platform-  
510 neutral text-based language called eXtensible Access Control Markup Language (XACML)  
511 Version 3.0, which has been standardized by OASIS. The other is Next Generation Access  
512 Control (NGAC), whose data structure and operations have been standardized under INCITS  
513 565-2020 [11] – Information technology – Next Generation Access Control. This standardization  
514 includes the APIs of functional components (i.e., PEP, PDP, RAP), allowing for the  
515 interoperability of these components from different sources. Further, the PEP interface is  
516 common for enforcing policies over both application requests and policy administration requests.  
517 The biggest advantage of NGAC is the use of linear time algorithms for computing access  
518 control decisions and performing policy reviews (i.e., determining the set of resources that a user  
519 can access, determining the set of users that can access a resource) [12,13].

520 The functional architectures for these two representational structures are given in Figures 3.1 and  
521 3.2, respectively. A brief description of the modules in these two functional architectures is as  
522 follows:


- 523 • Policy Decision Point (PDP) – This is the core module of the ABAC functional

- 524 architecture that computes decisions to permit or deny user access requests for  
525 performing actions on resources. Requests are received from and responses sent to a  
526 module called Policy Enforcement Point (PEP) in both representations.
- 527 • Policy Enforcement Point (PEP) – This is a module that is part of the application’s  
528 platform and tightly integrated with the application. It is designed to intercept all access  
529 requests that emanate from the application in both representations.
  - 530 • Policy Information Point (PIP)
    - 531 a. In the XACML representation, this is a module that contains the database of attributes  
532 and their associated values for all application-relevant objects or resources. The  
533 information here is used to extract the attributes and associated values for users and  
534 resources found in the access request to find the applicable target policies in the PRP  
535 (described below).
    - 536 b. In the NGAC representation, this is a repository of association relations of the form  
537 (u-ai, op-i, o-ai) for a pc-i, where u-ai and o-ai are attribute values associated with a  
538 user and object (resource), respectively. op-i denotes a set of allowed operations and  
539 pc-i the governing policy classes. To minimize the set of association relations in the  
540 authorization database (e.g., having triples to represent every user and every object in  
541 the application), containment relations of the form (U < u-ai) are used to show the  
542 members of the user group and object group represented in the association relations.  
543 In addition, the same set of containment relations are used to denote the applicable  
544 policies for each object as well (O < pc-i).
  - 545 • Policy Retrieval Point (PRP) – In the XACML representation, this is the module that is  
546 the repository for authorization policies expressed as logical formulas involving  
547 predicates on attribute values. The policy representation also contains the target resources  
548 that are covered by the policy. The resources requested in the access request are matched  
549 to these targets to retrieve the applicable policies by the PDP when computing decisions  
550 for those requests. This module is not part of the functional architecture in the NGAC  
551 representation.
  - 552 • Attribute Administration Point (AAP) – This is the interface for administering attributes  
553 stored in PIP in the XACML representation. This module is not necessary in NGAC  
554 representation since its association relations express the access rights on objects  
555 instantiated using attribute values.
  - 556 • Policy Administration Point (PAP) – This is the interface for administering policies  
557 stored in PRP.

558

559

560



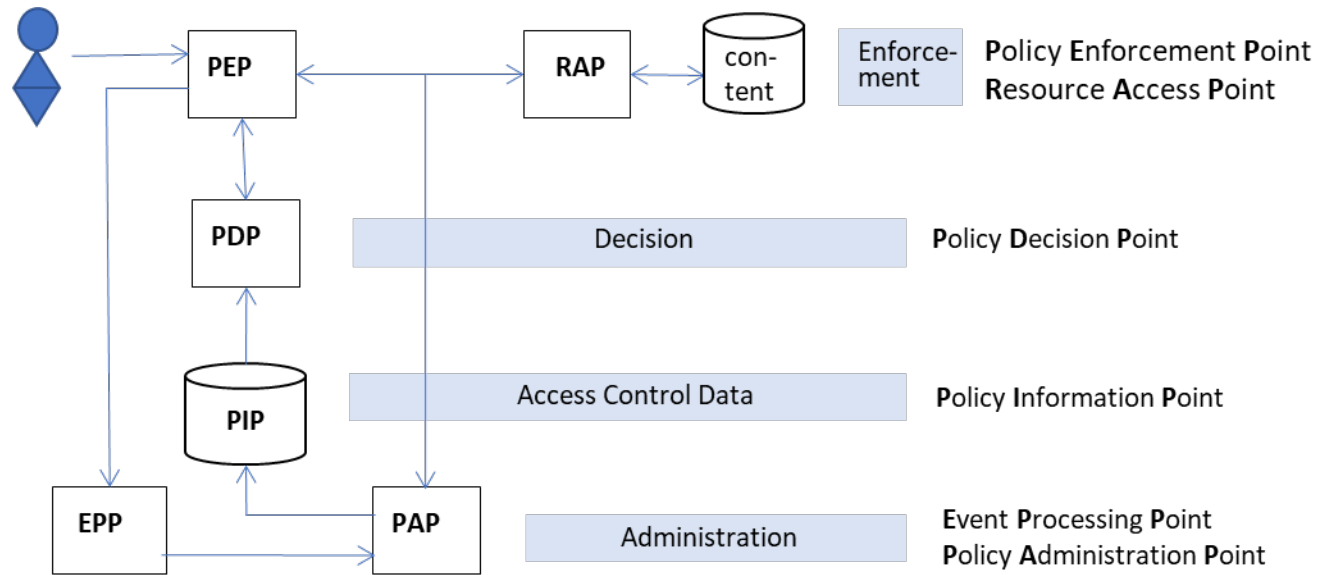
561

562

Figure 3.1 ABAC Functional Architecture based on XACML Representation

563

564



565

566

Figure 3.2 ABAC Functional Architecture based on NGAC Representation

567

568

**569 3.1 ABAC Deployment for Microservices-based Applications Using Service Mesh**

570 In the context of a microservices-based application using service mesh, an ABAC deployment  
571 for access control can take the following forms:

- 572 ● The proxies (e.g., Ingress, sidecar, and Egress) play the role of PEPs since they intercept all  
573 requests that emanate from each client, user, service, or external service.
- 574 ● The PEPs can provide either an ALLOW/DENY verdict or a list of allowable objects.
- 575 ● The enforcement function can be provided either natively (using local configuration  
576 structures, such as ACLs) or using proxy extensions that call an external authorization server  
577 to obtain one or more of the data in the previous bullet.
- 578 ● The assurance mechanisms in the service mesh (e.g., certificate-based authentication, secure  
579 session, non-bypassability, execution isolation) can be leveraged to deploy a high assurance  
580 authorization framework.

581

## 582 **4 Authentication and Authorization Policy Configuration in Service Mesh**

583 Fine-grained access control for microservices can be enforced through the configuration of  
584 authentication and access control policies. These policies are defined in the control plane of the  
585 service mesh, mapped into low-level configurations, and pushed into the sidecar proxies that  
586 form the data plane of the service mesh. The configurations enable the proxies to enforce the  
587 policies at application runtime (or request time), thus making the proxies act as Policy  
588 Enforcement Points (PEPs). As stated in the introduction, the objective of this document is to  
589 provide guidance for the deployment of an authentication and authorization framework that is  
590 external to the application, agnostic to the platform hosting the application and the service mesh  
591 product that implements the application infrastructure. However, Kubernetes is used as the  
592 reference application platform and Istio as the service mesh infrastructure platform to provide  
593 concrete examples of the concepts and to enable us to make specific recommendations with more  
594 clarity and specificity.

### 595 **4.1 Hosting Platform Configuration**

596 The generic host platform configuration data for microservices-based applications using service  
597 mesh that are, at the minimum, needed for authentication and authorization policy configuration  
598 are:

- 599 • Metadata, like service name and the sets of instances of that service
- 600 • Runtime data, such as services's protocols and ports
- 601 • Namespaces that provide logical isolation boundaries for sets of services
- 602 • Unique runtime identities for each service

603 In the reference hosting platform Kubernetes, this is realized as:

- 604 • Service resource, which declares a service's name, protocol (e.g. TCP), and ports (e.g.  
605 9080)
- 606 • Deployment resource, which declares deployments of pods that implement that service,  
607 including metadata such as labels and version
- 608 • Namespace construct and RBAC for managing how users are allowed to publish  
609 configuration into namespaces
- 610 • Service Accounts, which are identities unique to each namespace bound to individual  
611 services

### 612 **4.2 Service Mesh Configuration**

613 The installation of any service mesh involves the following components:

- 614 • Ingress Gateway, which is the first point of entry into the microservices-based application.  
615 This gateway specification includes names, ports, and routes that the application client  
616 must take to access the application.

- 617 ● Egress Gateway for the application to call outside services or applications. Egress  
618 gateways are optional since a sidecar proxy can act as an egress proxy for the purposes of  
619 policy without deploying an egress gateway.
- 620 ● Injection of sidecar proxies (in the form of containers). The consequence of this is that  
621 each of the application's deployments in the platform will now have two containers—the  
622 original microservice container plus the mesh's sidecar proxy. These sidecar proxies  
623 enforce authentication and authorization policies during application runtime, thus acting  
624 as Policy Enforcement Points (PEPs). In addition, proxies should emit metrics and logs to  
625 enable continuous monitoring of the system; this can be used to ensure policies are in  
626 place and are being enforced.
- 627 ● A Certificate Authority (CA) module is needed to handle certificate requests from sidecar  
628 proxies, which need a runtime identity presented as an X.509 certificate. This CA  
629 generates, distributes, and manages keys and certificates used by the mesh and enables the  
630 mesh to perform automatic certificate rotation. A CRL or OCSP feature is also required to  
631 support certificate validation.
- 632 ● A control plane module in the service mesh that monitors configuration data in the hosting  
633 platform, encodes policies and distributes those policies in the form of configuration to  
634 various proxies in the mesh (e.g., Ingress, sidecar, and Egress).

635 In the context of the reference service mesh platform Istio, to facilitate route specification to the  
636 entry service of the application in the Ingress Gateway, a virtual service is defined that specifies  
637 the path and hosts making up the virtual service and the first entry service/port to which the  
638 gateway must route the incoming request from an application client [14].

639 **SMC-SR-1:** *The signing certificate used by the mesh's CA module should be rooted in the*  
640 *organization's existing PKI to allow for auditability, rotation, and revocation.*

641 Some service meshes come with the ability to encrypt traffic using a self-signed certificate; such  
642 a certificate should not be used in secure deployments.

643 **SMC-SR-2:** *Communication between the service mesh control plane and the hosting platform's*  
644 *configuration server must be authenticated and authorized.*

645 In this reference platform, authentication is typically achieved by the Kubernetes API Server (the  
646 configuration server) with simple TLS. Authentication of the client is based on the pod's service  
647 account credential. Authorization for the client to receive platform information from the API  
648 Server is enforced by Kubernetes RBAC.

### 649 **4.3 Higher-level Security Configuration Parameters**

650 Since the component microservices of our application are generally implemented as containers,  
651 the following higher-level security configuration parameters should be set. In the reference  
652 hosting platform Kubernetes, containers are implemented in pods, which contain a microservice  
653 container as well as a sidecar container. These higher-level security configurations are set through  
654 flags that come under the banner of pod security policies. The recommendations for these flag

655 values are numbered using the acronym HLC-SR-X, where HL stands for higher-level  
656 configuration, SR stands for security recommendation, and X is the sequence number. They  
657 include but are not limited to the following [5]:  
658

659 **HLC-SR-1:** *Containers and applications should not be run as root (thus becoming privileged*  
660 *containers).*

661 In Kubernetes, the configuration setting for this is to set the value TRUE for  
662 “`MustRunAsNonRoot`” flag.  
663

664 **HLC-SR-2:** *Host path volumes should not be used as they create tight coupling between the*  
665 *container and the node on which it is hosted, constraining the migration and flexible resource*  
666 *scheduling process.*

667 In Kubernetes, the configuration setting for this is to set the value of TRUE to  
668 “`readOnlyRootFilesystem`” flag.  
669

670 **HLC-SR-3:** *Configure the container file system as read-only by default for all applications,*  
671 *overriding only when the underlying application (e.g., database) must write to disk.*  
672

673 **HLC-SR-4:** *Explicitly prevent privilege escalation for containers.*

674 In Kubernetes, this is achieved by setting the value FALSE for the  
675 “`allowPrivilegeEscalation`” flag.

#### 676 **4.4 Authentication Policies**

677 Authentication policies specify the process for validating identities. The integrity of this process  
678 and its strength determines the integrity of the authorization process since the latter depends upon  
679 the strength of the authenticated identity. There are two types of identity needed in a  
680 microservices-based application:

- 681 • Microservices or workload identity
- 682 • End-user identity

683 Service (microservice) identity is critical for the following reasons:

- 684 • It enables the client to verify that the server to which it is communicating (server identity  
685 validated using the certificate it carries) is authorized to run the service. This assurance  
686 has to be provided by a secure naming service that maps the server identity to the service  
687 identity. In any orchestration platform (including Kubernetes), services can be moved  
688 around the nodes (server) for load balancing and service availability reasons. It is the  
689 responsibility of the control plane of the service mesh to refresh this mapping information  
690 by interacting with the API that contains this configuration information (e.g., through API  
691 server in Kubernetes) and convey it to the sidecar proxy in the data plane of the service  
692 mesh.
- 693 • The service identity is the basis for the target service to select and enforce applicable  
694 authorization policies.

#### 695 4.4.1 Specifying Authentication Policies

696 Associated with these identities are the corresponding authentication processes that the service  
697 meshes have to support. They are:

- 698 ● Service-level authentication or peer authentication using service identity
- 699 ● End user authentication or request authentication using end user credentials

700 It is assumed that the reference hosting platform has been configured with the high-level  
701 requirements outlined in Section 4.1. It is also assumed that the reference service mesh platform  
702 has been installed and configured with the initial requirements outlined in Section 4.2.

#### 703 4.4.2 Service-level Authentication

704 Service-level authentication is the mutual authentication of the communicating services and setup  
705 of a secure TLS session. Enabling this requires the capability to define a policy object which  
706 should meet the following requirements:

707 **AUN-SR-1:** *A policy object relating to service-level authentication should be defined that requires*  
708 *that mTLS be used for communication. The policy object should be expressive enough to be defined*  
709 *at various levels (given below) with features for overrides at the lower levels or inheritance of the*  
710 *requirement specified at the higher levels. The following are the minimum required levels [6]:*

- 711 a. Global level or the service mesh level
- 712 b. Namespace level
- 713 c. Workload or microservices level – used for applying authentication and authorization  
714 policies for a subset of traffic to a subset of resources (e.g., particular microservices, hosts  
715 or ports)
- 716 d. Port level, taking into account that certain traffic is designed for communicating through  
717 designated ports

718 This form of authentication also requires the assignment of a strong identity to each service and  
719 the authenticating of that identity by mapping it to the server identity (where the service is  
720 hosted) that digitally signed in a special digital authentication certificate (SPIFFE). To provide  
721 assurance that the server whose identity is found in the SPIFFE certificate is the one that is  
722 authorized to run the target service, the following requirement (also specified in SP 800-204A) is  
723 needed:

724 **AUN-SR-2:** *If the certificate used for mTLS carries server identity, then the service mesh should*  
725 *provide a secure naming service that maps the server identity to the microservice name that is*  
726 *provided by the secure discovery service or DNS. This requirement is needed to ensure that the*  
727 *server is the authorized location for the microservices and to protect against network hijacking.*

728 The information for mapping the server identity to a service is obtained by the control plane of  
729 the service mesh by accessing the configuration information from the platform that is hosting the



730 microservices-based application. In Kubernetes, the control plane of the service mesh obtains the  
731 mapping information through the API server module of the Kubernetes platform and populates  
732 that information in the secure naming service. Thus, the mutual certificate validation not only  
733 enables validation of the associated service identities of both the client and target services but  
734 also enables creation of a secure mutual TLS (mTLS) session. In Istio, the policy object for this  
735 type of authentication is called “peer authentication.”

#### 736 4.4.3 End User Authentication

737 For the mesh to authenticate end user credentials (EUC), the application must participate in some  
738 way. Client services that make the request should acquire and attach an appropriate credential to  
739 each request (e.g., a JWT) in the request header. End user authentication, or request  
740 authentication, is the process of validating the credentials of the end user making a request by  
741 extracting from the request’s metadata and authenticating them (locally or against an external  
742 server). For example, a common flow at many organizations is to exchange an external EUC, like  
743 an Oauth bearer token, at ingress for an internal credential that is encoded within a JSON Web  
744 Token (JWT). The JWT can be created by a custom authentication provider or standards-based  
745 OpenID Connect provider.

746 **EAUN-SR-1:** *A request authentication policy must, at the minimum, provide the following*  
747 *information:*

- 748 ● *Instructions for extracting the credential from the request*
- 749 ● *Instructions for validating the credential*

750 For a JWT, this might include:

- 751 ● Location (header name) of the JWT token that contains the user’s claims
- 752 ● How to extract the subject, claims, and issuers from the JWT
- 753 ● Public keys or the location for the key used for validating the JWT

#### 754 4.5 Authorization Policies

755 Authorization policies, just like their authentication counterparts, can be specified at the service  
756 level as well as the end user level. In addition, authorization policies are expressed based on  
757 constructs of an access control model and thus may vary based on the nature of the application  
758 and enterprise-level directives. Further, the location of the access control data may vary  
759 depending on the identity and access management infrastructure in the enterprise. These  
760 variations result in the following variables:

- 761 ● Two authorization levels – service level and end user level
- 762 ● Access control model used to express authorization policies
- 763 ● Location of the access control data in a centralized or external authorization server or  
764 carried as header data

765 The supported access control in the service mesh uses abstraction to group one or more policy  
766 components (described below in Section 4.5.1) for specifying either service-level or end user-  
767 level authorization policies. Since microservices-based applications are implemented as APIs  
768 (e.g., RESTful API), authorization policy components described using key/value pairs will have  
769 attributes pertaining to an API, including the associated network protocols. The types of  
770 authorization policies are:

- 771 ● Service-level authorization policies
- 772 ● End user-level authorization policies
- 773 ● Model-based authorization policies

#### 774 4.5.1 Service-level Authorization Policies

775 Service-level authorization policies are defined using a policy object that provides positive or  
776 negative permission (authorization) with the following policy components:

- 777 a. The scope of the policy can span all applications at the service mesh level, namespace  
778 level, or one or more designated applications (microservice level).
- 779 b. The permissions or operations can be restricted to one or more designated methods of a  
780 given service (e.g., an “HTTP GET method on the ‘/details’ path of an application named  
781 PRODUCT-CATALOG”) or to designated ports through which an application can be  
782 accessed.
- 783 c. Conditions under which access can take place (e.g., possession of a token) are specified.
- 784 d. Sources allowed access are specified at the namespace or a particular service level (in  
785 terms of the service’s runtime identity).

786 **AUZ-SR-1:** *A policy object describing service-to-service access should be in place for all*  
787 *services in the mesh. At a minimum, these policies should permit access at the namespace level*  
788 *(e.g., “services in namespace A can call services in namespace B”).*

789 Ideally, policies should describe the minimum access required for application functionality (e.g.,  
790 “service ‘foo’ in namespace A can perform ‘GET /bar’ on service ‘bar’ in namespace B”).

#### 791 4.5.2 End-user Level Authorization Policies

792 Given an authentication policy like Section 4.4.3, a sidecar in the mesh can extract a principal  
793 from the request to perform authorization on. Further, the sidecar typically has additional context  
794 about the request, including the resource being accessed (e.g., the path in an HTTP/REST API)  
795 and the action being taken (e.g., the HTTP verb – GET, PUT, etc. – in the request to that API).  
796 This gives the sidecar enough information to act as a policy enforcement point and call a policy  
797 decision point.

798 This is the most common case, especially for organizations with traditional IAM systems that  
799 exist as an external service, often called by an SDK. To handle this case, a service mesh’s sidecar  
800 proxy will typically support calling external services to render an authentication and authorization

801 verdict. For example, the reference implementation Istio supports this via Envoy's (i.e., the  
802 sidecar proxy) external authorization service [15].

803 **EUAZ-SR-1:** *When a sidecar communicates with an authentication or authorization system, that*  
804 *communication should be secured with the mesh's built-in service-to-service authentication and*  
805 *authorization capabilities.*

806 Logs and metrics exported by the sidecar can be used to prove that authentication and  
807 authorization was performed by the sidecar on behalf of the application.

808 End user authorization is not applied to the decision endpoint of the external authorization (PDP)  
809 service since the service is the principal making the call. It also avoids needing a default policy  
810 that allows all users to call the decision endpoint of the PDP. End user authorization should be  
811 applied to the PAP and other administrative endpoints of the authorization system, and that can  
812 be facilitated by the mesh.

813 However, there is another case that is common enough to address in which an external  
814 authorization system is not required. Making a network call to an authorization service for every  
815 hop in a service chain can be expensive and cause centralized failures. To mitigate these  
816 problems, many organizations will exchange end user credentials at ingress for an internal,  
817 trusted, authenticatable credential that conveys not just the user's principal but also that user's  
818 capabilities in the system. A JSON Web Token (JWT) is frequently used for this because it is  
819 locally authenticatable and conveys the user's principal (the JWT's subject), the issuer of the  
820 JWT (issuer), and arbitrary claims that the organization can control (e.g., to use for access  
821 control).

822 Performing end user authorization based on a JWT is common enough that it is built directly into  
823 Envoy, the sidecar proxy of the reference mesh, Istio. Envoy can be configured with a filter [16]  
824 that will process requests in two steps:

- 825 a. JWT token verification involves extracting the token from the request header, verifying  
826 whether issuers and audiences are allowed, fetching the public key, and verifying the  
827 digital signature on the token.
- 828 b. Match the resources in the request to the claims in the token to determine whether the end  
829 user should be allowed access to the requested resources or denied.

830 Envoy's JWT filter act as the PDP, making the access decision entirely locally. This requires that  
831 policy documents be small enough to reside on an individual sidecar proxy. Although a full  
832 ABAC is ideal for handling resource-level policies, the JWT filter is valuable as a stepping stone  
833 from a traditional system that only performs access control on the edge to a zero trust system that  
834 performs authentication and authorization at each service.

```

apiVersion: security.istio.io/v1beta1
kind: AuthorizationPolicy
metadata:
  name: backend
  namespace: product
spec:
  action: ALLOW
  rules:
  - from:
    - source:
      principals: ["cluster.local/ns/product/sa/frontend"]
    to:
    - operation:
      methods: ["GET"]
      paths: ["/info*"]
    - operation:
      methods: ["POST"]
      paths: ["/data"]
  when:
  - key: request.auth.claims[iss]
    values: ["accounts.google.com"]

```

**Figure 4.1 – An example Istio authorization policy**

This allows the front end to call specific methods on the backend only if the request has an EUC attached issued by “accounts.google.com.”

835

836 **EUAC-SR-2:** *All application traffic should carry end-user credentials, and there should be a*  
 837 *policy in the mesh enforcing that credentials are present.*

838 We recommend this even if the application is enforcing authentication and authorization  
 839 independently of the mesh, because these organization-wide controls allow functionality like  
 840 audit to be built on top of the mesh at lower cost to central teams responsible for compliance and  
 841 controls.

#### 842 **4.5.3 Model-based Authorization Policies**

843 The service-level authorization policies and a use case of end-user authorization policies that uses  
 844 JWT are natively implemented in the proxies. Since these cannot be used for resource-level  
 845 authorization policies, we need to support model-based authorization policies as well. As already  
 846 alluded to in section 4.5.2, this requires a call from the proxy to an external authorization server  
 847 which holds the model-based authorization engine to obtain an access decision.

848 The service principals in these model-based policies are identities (e.g., ServiceAccount) that are  
849 provided by the underlying application orchestration platform (e.g., Kubernetes) and is the same  
850 that are used by authorization policies natively supported in the proxies. The user principals are  
851 usually obtained from the JWT token. The popular access control models in the external  
852 authorization servers are either RBAC or ABAC.4.6 Authorization Policy Elements

## 853 **4.6 Authorization Policy Elements**

854 The authorization policies that can be specified in a service mesh may consist of the following  
855 elements:

- 856 ● The policy types – Positive (ALLOW) or Negative (DENY)
- 857 ● The policy target or authorization scope – The namespace, a particular service  
858 (application name), and version
- 859 ● The policy sources – Covers the set of authorized services
- 860 ● The policy operations – Specifies the operations on the target resources that are covered  
861 under the policy
- 862 ● The policy conditions – The metadata associated with the request that must be met for the  
863 application or invocation of the policy

### 864 **4.6.1 Policy Types**

865 Positive and negative policies are specified in order to set precedence relationships (e.g., DENY  
866 overrides, ALLOW, etc.). They are also used for situations that allow one type of policy for all  
867 services under a group and to specify exceptions (e.g., have an ALLOW policy for all services in  
868 a namespace but a DENY policy for a specified service)

### 869 **4.6.2 Policy Target or Authorization Scope**

870 This refers to the target resources in terms of a set of services, versions, and the namespaces  
871 under which the services are located. The service can be specified in the following ways:

872 Using path: The location of the target resource is specified using paths (e.g., for resources  
873 accessed using HTTP or gRPC protocols). The list of paths to be included in the authorization  
874 policy scope and paths that need to be excluded can be defined. Both of these sub-elements of the  
875 policy target component (i.e., the list of paths to be included and the list of paths to be excluded)  
876 are optional.

877 Using host name: In some instances, the target resources are specified using the host sub-element.  
878 The list of hosts to be included in the authorization policy scope as well as those hosts that need  
879 to be excluded can be defined. Both of these sub-elements of the policy target component (i.e.,  
880 list of hosts to be included and the list of hosts to be excluded) are optional.

881 Using network ports: The network port through which the target resource (the service) is accessed  
882 is often specified using the port sub-element. The list of ports to be included in the authorization

883 policy scope as well as those ports that need to be excluded can be defined. Both of these sub-  
884 elements of the policy target component (i.e., list of ports to be included and the list of ports to be  
885 excluded) are optional.

#### 886 **4.6.3 Policy Sources**

887 The policy sources are the set of services that are authorized to operate on the set of resources  
888 specified under the policy target (specified using name, path, host name, and ports). The policy  
889 sources are usually specified using a service account or name (called principal), all services in a  
890 particular logical group (e.g., namespace), or all services that are accessed from a group of  
891 network locations (e.g., IP blocks). Both included and excluded principals, namespaces, and IP  
892 blocks can be specified in some implementations.

#### 893 **4.6.4 Policy Operations**

894 The set of operations depends on the way the application is implemented. If the application is  
895 implemented as a REST API, the following are the common operations (also called HTTP verbs  
896 or HTTP methods):

897       POST: This is equivalent to creating a resource.

898       GET: This is equivalent to reading the contents of the resource.

899       PUT: This is equivalent to updating the resource by replacing.

900       PATCH: This is equivalent to updating the resource by modifying.

901       DELETE: This is equivalent to deleting the resource.

902       OPTIONS:

903       HEAD:

904 If the resource is accessed using gRPC instead of a RESTful protocol, there is only one operation  
905 or method: “POST.” The authorization policy definition may also have a feature to specify the list  
906 of operations (methods) to be excluded. Both policy sub-elements—one to specify the operations  
907 to be included in the authorization policy scope and the other to be excluded—are optional.

#### 908 **4.6.5 Policy Conditions**

909 Policy conditions specify the constraints in the form of a key-value pair for the metadata  
910 associated with the request. This metadata may cover the following:

911 *Metadata associated with the source:* Some of the metadata (e.g., service account name,  
912 namespace, and IP blocks) are specified as part of the policy source specification itself. In  
913 addition, it is possible to list IP addresses in CIDR format of the policy sources.

914 *Metadata associated with the request:* In this type of metadata, the parameters or attributes that  
915 pertain to a specific request can be specified. These parameters can include an audience that can  
916 present the authentication information expressed in the form of a URL (only applicable to HTTP  
917 protocol-based requests), a specific end user identifier associated with the audience that can  
918 present the authentication credentials, or the claim name that is carried in the token presented by  
919 the presenter. In addition, parameters that pertain to the user-agent (e.g., browser name) can also  
920 be specified for HTTP protocol-based requests.

921 *Metadata associated with the destination:* The range of allowable IP addresses can be specified in  
922 CIDR format as well as the associated list of ports.

923

924

## 925 **5 ABAC Deployment for Service Mesh**

926 The last chapter introduced three different types of authorization policies including two use cases  
927 for end-user level authorization policies. This chapter, we will leverage those architectural choices  
928 to describe an ABAC-based authorization framework in the service mesh:

- 929 ● Security assurance for authorization framework enforcement
- 930 ● Supporting infrastructure for authorization requests
- 931 ● Advantages of ABAC Authorization framework for Service Mesh
- 932 ● Enforcement alternatives in Proxies

### 933 **5.1 Security Assurance for Authorization Framework Enforcement**

934 The authorization policy enforcement mechanism implemented in the service mesh for a  
935 microservices-based application must satisfy the three requirements of a reference monitor  
936 concept. It must be 1) non-bypassable, 2) protected from modification, and 3) verified and tested  
937 to be correct. These three requirements can be ensured by the following:

- 938 ● Every request from a client to the microservices-based application, from one service to  
939 another (inter-services call), and from a microservice to an external application is  
940 intercepted by the ingress gateway, sidecar proxy, and egress proxy, respectively, and  
941 these policy enforcement points (PEPs) are non-bypassable.
- 942 ● The policy enforcement modules are independent executables that are decoupled from the  
943 application logic and cannot be modified.
- 944 ● Their outcome can be independently verified and tested through both shadow operations  
945 and live production requests.

946 In short, a proxy running in the data plane of the service mesh is the reference monitor with  
947 respect to authorization enforcement. The authorization policy engine (e.g., NGAC-based ABAC  
948 policy engine) implemented as a container executing either natively in the proxy memory space  
949 or callable from a corresponding filter module in the proxy runs as a separate process that does  
950 not share any memory space with the calling application. Hence, it satisfies the requirement of a  
951 security kernel.

### 952 **5.2 Supporting Infrastructure for ABAC Authorization Framework**

953 We will now look at the basic building blocks of the supporting infrastructure for service-to-  
954 service and end-user+service-to service requests.

#### 955 **5.2.1 Service-to-Service Request (SVC-SVC) – Supporting Infrastructure**

956 The policy object used for authorizing this type of request was described in Section 4.5.1.  
957 Service-to-service requests must be authorized based on the identity of the calling and called  
958 services. The trusted document that carries the identity of the service is an X.509 certificate



959 issued by one of the control plane components of the service mesh after verifying whether the  
960 requested identity is valid for the microservice by consulting an identity registry. The proxy  
961 communicates with this control plane component through a local agent, obtains a certificate, and  
962 sends it to the proxy, which then performs the certificate validation process on behalf of the  
963 calling service or client during each service request. The identity is encoded as URI and carried in  
964 a certificate's SAN (subject alternate name) field. It must be mentioned that the certificates that  
965 carry service account identities are short-lived certificates (rotated every hour or few hours)  
966 rather than the conventional HTTPS TLS terminating certificates whose validity lasts for several  
967 months.

## 968 **5.2.2 End User + Service-to-Service Request (EU+SVC-SVC) – Supporting Infrastructure**

969 The policy object used for authorizing this type of request was described in Section 4.5.2. This  
970 request type requires the verification of two identities: the calling user identity and the service  
971 identity. As described in the previous section, the service mesh provides the feature to perform  
972 authorization based on service identities. Since this is a standard feature, no extra components  
973 need to be built in the service mesh infrastructure for this type of authorization. However, when  
974 end user identities are introduced for authorization, the authorization framework should be tightly  
975 integrated with the following components of the architecture:

- 976 ● The services orchestration control plane for obtaining application object attributes as well  
977 as attributes of the registered application users (which includes user credentials), thus  
978 playing the role of Policy Information Point (PIP) in ABAC-based authorization
- 979 ● A service mesh control plane for obtaining tokens that encode the claims based on the  
980 authorization decision
- 981 ● A service mesh data plane in the service proxy for making calls to the authorization  
982 engine (which is just another service), obtaining the authorization decision, enforcing the  
983 service-to-service authorization policies, making calls to the service mesh control plane  
984 for authorization tokens (e.g., Java Web Tokens [JWT]), and attaching the tokens to the  
985 service request.

986 The advantage of an EU+SVC-SVC request processing scheme is that authorizations at a finer  
987 level of granularity than the method level can be specified, and conformant claims can be  
988 included in the authorization token.

989 A disadvantage is that there is overhead involved in enforcing two layers of authorization—one  
990 layer based on policies specified for SVC-SVC requests and a second layer based on EU+SVC-  
991 SVC requests. Access control processing logic based on the second layer involves multiple calls  
992 by service proxy, such as (a) a call to the authorization engine service to obtain the access  
993 decision after obtaining the user attributes (including user credentials) and application object  
994 attributes from the orchestration system, (b) obtaining the authorization token from the service  
995 mesh control plane based on the access decision, and (c) including the authorization token along  
996 with service request.

### 997 **5.3 Advantages of ABAC Authorization Framework for Service Mesh**

998 We provide here the justification for the various building blocks of the architecture for our  
999 authorization framework – the service mesh, the NGAC-based ABAC model etc. We also  
1000 highlight the scalability and flexibility of certain components such as proxy APIs, NGAC  
1001 authorization engine etc.

- 1002 a. A service mesh is the right architecture for the enforcement of authorization policies since  
1003 the components involved are moved out of the application and executed in a space where  
1004 they can form a security kernel that can be vetted.
- 1005 b. Both types of authorization requests (i.e., SVC-SVC and EU-SVC-SVC) can be handled  
1006 by a runtime infrastructure that involves the coupling of orchestration platform control  
1007 plane, service mesh control plane, and mesh data plane to the access control engine.
- 1008 c. The extensible API of the proxy can be used to integrate any authorization engine using  
1009 the appropriate type of access control model. ABAC has been found to be one of the most  
1010 flexible, scalable access control models because of its ability to incorporate any number  
1011 and type of attributes associated with the subject, object, and environment.
- 1012 d. Performance requirements for the authorization engine are met due to the linear time  
1013 processing speed of the graph-based, NGAC-based ABAC model.
- 1014 e. The flexibility outlined in (c) can be leveraged to incorporate models for both application  
1015 and data protection. Enabling data protection models such as NDAC can be part of the  
1016 authorization server.

### 1017 **5.4 Enforcement Alternatives in Proxies**

1018 Authorization can be enforced through a native structure (e.g., authorization policy) supported in  
1019 the particular version of the service mesh or using calls to an external authorization server. The  
1020 external authorization server can use any access control model and any representation of policy  
1021 expressions (logical rules or acyclic graph representations), but the mediation of a request coming  
1022 into the proxy can be performed in the following ways:

- 1023 a. Each request is passed on to the external authorization server through the external  
1024 authorization filter in the proxy, and the response from the authorization server is used for  
1025 request mediation in the form of ALLOW or DENY.
- 1026 b. Prestored ACLs can be used in the proxy itself, generated by calls to the authorization  
1027 server. If the authorization server uses an enterprise-wide access control model, an  
1028 administrative API may be needed that will perform the function of mapping the  
1029 enterprise resources to resources, users, and groups pertaining to the service served by its  
1030 proxy to generate ACLs that are customized for the service.

1031

**6 Summary and Conclusions**

1033 Deployment guidance has been provided for an ABAC-based authorization framework for  
1034 securing microservices-based applications using a service mesh. Background information in  
1035 terms of authentication and authorization policies natively supported in proxies of the service  
1036 mesh are discussed. For supporting any authentication and authorization framework in the mesh,  
1037 the pre-requisites in the form of hosting platform configuration data, the service mesh  
1038 configuration and some higher-level security configuration parameters for orchestration of  
1039 component microservices (when implemented as containers) are outlined.

1040 The description of the ABAC deployment in the service mesh includes the requirements for  
1041 security assurance, supporting infrastructure, advantages of ABAC authorization framework and  
1042 enforcement alternatives in proxies.

1043

1044 **References**

- 1045 [1] Rose S, Borchert O, Mitchell S, Connelly S (2020) Zero Trust Architecture. (National  
1046 Institute of Standards and Technology, Gaithersburg, MD), NIST Special Publication (SP)  
1047 800-207. <https://doi.org/10.6028/NIST.SP.800-207>
- 1048 [2] Red Hat (2021) *What is CI/CD?* Available at  
1049 [https://www.redhat.com/en/topics/devops/what-is-ci-  
1050 cd#:~:text=CI%2FCD%20is%20a%20method,continuous%20delivery%2C%20and%20c  
1051 ontinuous%20deployment](https://www.redhat.com/en/topics/devops/what-is-ci-cd#:~:text=CI%2FCD%20is%20a%20method,continuous%20delivery%2C%20and%20c,ontinuous%20deployment)
- 1052 [3] Red Hat (2021) *What is DevSecOps?* Available at  
1053 <https://www.redhat.com/en/topics/devops/what-is-devsecops>
- 1054 [4] Chandramouli R (2019) Security Strategies for Microservices-based Application Systems.  
1055 (National Institute of Standards and Technology, Gaithersburg, MD), NIST Special  
1056 Publication (SP) 800-204. <https://doi.org/10.6028/NIST.SP.800-204>
- 1057 [5] Hu VC, Ferraiolo DF, Chandramouli R, Kuhn DR (2018) Attribute-Based Access Control  
1058 (Artech House, Boston USA).
- 1059 [6] Chandramouli R, Butcher Z (2020) Building Secure Microservices-based Applications  
1060 using Service-Mesh Architecture. (National Institute of Standards and Technology,  
1061 Gaithersburg, MD), NIST Special Publication (SP) 800-204A.  
1062 <https://doi.org/10.6028/NIST.SP.800-204A>
- 1063 [7] McEvoy E (2019) *Cordanetes: Combining Corda and Kubernetes* (Medium.com).  
1064 Available at <https://medium.com/corda/combining-corda-and-kubernetes-4e2ba54494c7>
- 1065 [8] Ramakani A (2020) *Kong API Gateway – From Zero to Production* (Medium.com).  
1066 Available at <https://medium.com/swlh/kong-api-gateway-zero-to-production-5b8431495ee>
- 1067
- 1068 [9] Agarwal G (2020) *How to Manage Microservices on Kubernetes With Istio*  
1069 (Medium.com). Available at [https://medium.com/better-programming/how-to-manage-  
1070 microservices-on-kubernetes-with-istio-c25e97a60a59](https://medium.com/better-programming/how-to-manage-microservices-on-kubernetes-with-istio-c25e97a60a59)
- 1071
- 1072 [10] Agarwal G (2020) *How Istio Works Behind the Scenes on Kubernetes* (Medium.com).  
1073 Available at [https://medium.com/better-programming/how-istio-works-behind-the-scenes-  
1074 on-kubernetes-aeb8003f2cb5](https://medium.com/better-programming/how-istio-works-behind-the-scenes-on-kubernetes-aeb8003f2cb5)
- 1075
- 1076 [11] InterNational Committee for Information Technology Standards (2020) *INCITS 565-2020*  
1077 *- Information technology - Next Generation Access Control* (INCITS, Washington, DC).  
1078 Available at  
1079 [https://standards.incits.org/apps/group\\_public/project/details.php?project\\_id=2328](https://standards.incits.org/apps/group_public/project/details.php?project_id=2328)

- 1080 [12] Mell P, Shook J, Harang R, Gavril S (2017) Linear Time Algorithms to Restrict Insider  
1081 Access using Multi-Policy Access Control Systems. *Journal of Wireless Mobile Networks,*  
1082 *Ubiquitous Computing, and Dependable Applications* 8(1):4-25.  
1083 <https://doi.org/10.22667/JOWUA.2017.03.31.004>
- 1084 [13] Ferraiolo D, Chandramouli R, Hu V, Kuhn R (2016) A Comparison of Attribute Based  
1085 Access Control (ABAC) Standards for Data Service Applications: Extensible Access  
1086 Control Markup Language (XACML) and Next Generation Access Control (NGAC)  
1087 (National Institute of Standards and Technology, Gaithersburg, MD), NIST Special  
1088 Publication (SP) 800-204A. <https://doi.org/10.6028/NIST.SP.800-178>
- 1089 [14] Agarwal G (2020) *Enable Mutual TLS Authentication between your Kubernetes*  
1090 *Workloads Using Istio* (Medium.com). Available at [https://medium.com/better-](https://medium.com/better-programming/enable-mutual-tls-authentication-between-your-kubernetes-workloads-using-istio-65338c8adf82)  
1091 [programming/enable-mutual-tls-authentication-between-your-kubernetes-workloads-](https://medium.com/better-programming/enable-mutual-tls-authentication-between-your-kubernetes-workloads-using-istio-65338c8adf82)  
1092 [using-istio-65338c8adf82](https://medium.com/better-programming/enable-mutual-tls-authentication-between-your-kubernetes-workloads-using-istio-65338c8adf82)
- 1093 [15] Envoy (2021) *External Authorization*. Available at  
1094 [https://www.envoyproxy.io/docs/envoy/v1.17.0/intro/arch\\_overview/security/ext\\_authz\\_fi](https://www.envoyproxy.io/docs/envoy/v1.17.0/intro/arch_overview/security/ext_authz_filter#arch-overview-ext-authz)  
1095 [lter#arch-overview-ext-authz](https://www.envoyproxy.io/docs/envoy/v1.17.0/intro/arch_overview/security/ext_authz_filter#arch-overview-ext-authz)
- 1096 [16] Envoy (2021) *JWT Authentication*. Available at  
1097 [https://www.envoyproxy.io/docs/envoy/v1.17.0/configuration/http/http\\_filters/jwt\\_authn\\_](https://www.envoyproxy.io/docs/envoy/v1.17.0/configuration/http/http_filters/jwt_authn_filter#config-http-filters-jwt-authn)  
1098 [filter#config-http-filters-jwt-authn](https://www.envoyproxy.io/docs/envoy/v1.17.0/configuration/http/http_filters/jwt_authn_filter#config-http-filters-jwt-authn)
- 1099
- 1100
- 1101
- 1102
- 1103